

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import rcParams
df=pd.read_csv('C:\\Users\\gayat\\Desktop\\Nalaiya Thiran Lab\\Assignment
4\\Mall_Customers.csv')
df.head()

```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```

#Univariate Analysis
var="Annual Income (k$)"
sns.distplot(df[var])
#Univariate Analysis
sns.displot(df.Age)
#Univariate Analysis
var="Spending Score (1-100)"
sns.distplot(df[var])
C:\Users\gayat\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a future
version. Please adapt your code to use either `displot` (a figure-level function with
similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
#Bivariate Analysis
sns.lineplot(df["Age"],df["Spending Score (1-100)"])
C:\Users\gayat\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variables as keyword args: x, y. From version
0.12, the only valid positional argument will be `data`, and passing other arguments
without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
#Bivariate Analysis
sns.lineplot(df["Annual Income (k$)"],df["Spending Score (1-100)"])
C:\Users\gayat\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variables as keyword args: x, y. From version
0.12, the only valid positional argument will be `data`, and passing other arguments
without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
df=df.drop(columns=['CustomerID'],axis=1)
#Multivariate Analysis
df.hist(figsize=(16,16))
array([[
],

```

```

[,
  ]], dtype=object)
df.describe()
Age    Annual Income (k$)  Spending Score (1-100)
count  200.000000         200.000000         200.000000
mean   38.850000         60.560000         50.200000
std    13.969007         26.264721         25.823522
min    18.000000         15.000000         1.000000
25%    28.750000         41.500000         34.750000
50%    36.000000         61.500000         50.000000
75%    49.000000         78.000000         73.000000
max    70.000000         137.000000        99.000000
df.isnull().any()
Gender                False
Age                   False
Annual Income (k$)    False
Spending Score (1-100) False
dtype: bool
sns.boxplot(df["Spending Score (1-100)"])
C:\Users\gayat\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12,
the only valid positional argument will be `data`, and passing other arguments
without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
sns.boxplot(df["Annual Income (k$)"])
C:\Users\gayat\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12,
the only valid positional argument will be `data`, and passing other arguments
without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
q1=df["Annual Income (k$)"].quantile(0.25)
q3=df["Annual Income (k$)"].quantile(0.75)
IQR=q3-q1
upper_limit= q3 + 1.5*IQR
df.median()
C:\Users\gayat\AppData\Local\Temp\ipykernel_11108\2508664500.py:5:
FutureWarning: Dropping of nuisance columns in DataFrame reductions (with
'numeric_only=None') is deprecated; in a future version this will raise TypeError.
Select only valid columns before calling the reduction.
  df.median()
Age                36.0
Annual Income (k$)  61.5
Spending Score (1-100)  50.0
dtype: float64
df["Annual Income (k$)"]=np.where(df["Annual Income
(k$)"]>upper_limit,61.5,df["Annual Income (k$)"])

```

```
sns.boxplot(df["Annual Income (k$)"])
```

C:\Users\gayat\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12,
the only valid positional argument will be `data`, and passing other arguments
without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
#Label encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
df.Gender=le.fit_transform(df.Gender)
df.head()
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	19	15.0	39
1	1	21	15.0	81
2	0	20	16.0	6
3	0	23	16.0	77
4	0	31	17.0	40

```
from sklearn.preprocessing import scale
df_scaled=pd.DataFrame(scale(df),columns=df.columns)
df_scaled.head()
```

	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1.128152	-1.424569	-1.788777	-0.434801
1	1.128152	-1.281035	-1.788777	1.195704
2	-0.886405	-1.352802	-1.748853	-1.715913
3	-0.886405	-1.137502	-1.748853	1.040418
4	-0.886405	-0.563369	-1.708930	-0.395980

```
new_df=df
from sklearn import cluster
error=[]
for i in range(1,11):
    kmeans=cluster.KMeans(n_clusters=i,init='k-means++',random_state=0)
    kmeans.fit(new_df)
    error.append(kmeans.inertia_)
error
```

C:\Users\gayat\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:1036:
UserWarning: KMeans is known to have a memory leak on Windows with MKL,
when there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.

```
warnings.warn(
[297063.675,
201152.1081841432,
139326.23321730684,
100349.31619915171,
71452.15398255126,
54455.93879921247,
48690.465943332725,
```

[illegible]

1.93048895 0.97405788 1.15063446 0.92113416 1.41427389 1.52766619
0.19941007 0.9844757 1.17210444 0.69753909 0.62148368 0.82357382
1.37326191 0.05857745 0.52962486 0.64161313 1.79687195 1.1983186
0.89771305 1.47795389 0.38921265 1.32502445 0.41131211 0.95727301
0.3490399 1.20691389 0.91885744 0.22494368]

Actual value 58 1

40 1
34 1
102 1
184 0
198 0
95 1
4 1
29 1
168 0
171 2
18 1
11 1
89 1
110 1
118 1
159 2
35 1
136 0
59 1
51 1
16 1
44 1
94 1
31 1
162 0
38 1
28 1
193 2
27 1
47 1
165 2
194 0
177 2
176 0
97 1
174 0
73 1
69 1
172 0

Name: kclus, dtype: int32

```

prediction=tree_model.predict(x_test)
print("Prediction",prediction)
print("Actual value",y_test)
Prediction [1.      1.      1.      1.      0.46666667 0.63333333
1.      1.      1.13333333 0.      2.      1.
1.03333333 1.      1.      1.      2.      1.
0.      1.      1.      1.      1.      1.
1.      0.      1.      1.      2.      1.
1.      2.      0.      1.93333333 0.      1.
0.      1.      1.      0.      ]
Actual value 58      1
40      1
34      1
102     1
184     0
198     0
95      1
4        1
29      1
168     0
171     2
18       1
11       1
89       1
110      1
118      1
159      2
35       1
136      0
59       1
51       1
16       1
44       1
94       1
31       1
162      0
38       1
28       1
193      2
27       1
47       1
165      2
194      0
177      2
176      0
97       1

```

```
174 0
73 1
69 1
172 0
Name: kclus, dtype: int32
from sklearn.metrics import mean_squared_error
print("Train: ", np.sqrt(mean_squared_error(y_train, model.predict(x_train))))
print("Test: ", np.sqrt(mean_squared_error(y_test, model.predict(x_test))))
Train: 0.39694877865218653
Test: 0.39489961809131696
print("Train: ", np.sqrt(mean_squared_error(y_train, tree_model.predict(x_train))))
print("Test: ", np.sqrt(mean_squared_error(y_test, tree_model.predict(x_test))))
Train: 0.06476453075908482
Test: 0.1267105187249881
y_pred = tree_model.predict(x_test)
fig = plt.figure(figsize=(10, 6))
plt.scatter(range(y_test.shape[0]), y_test, color='red', label='y_true')
plt.scatter(range(y_test.shape[0]), y_pred, color='blue', label='y_pred')
plt.legend()
plt.show()
```